

Lecture 14 : Multitasking


Slides adapted from Edward A. Lee

Outline

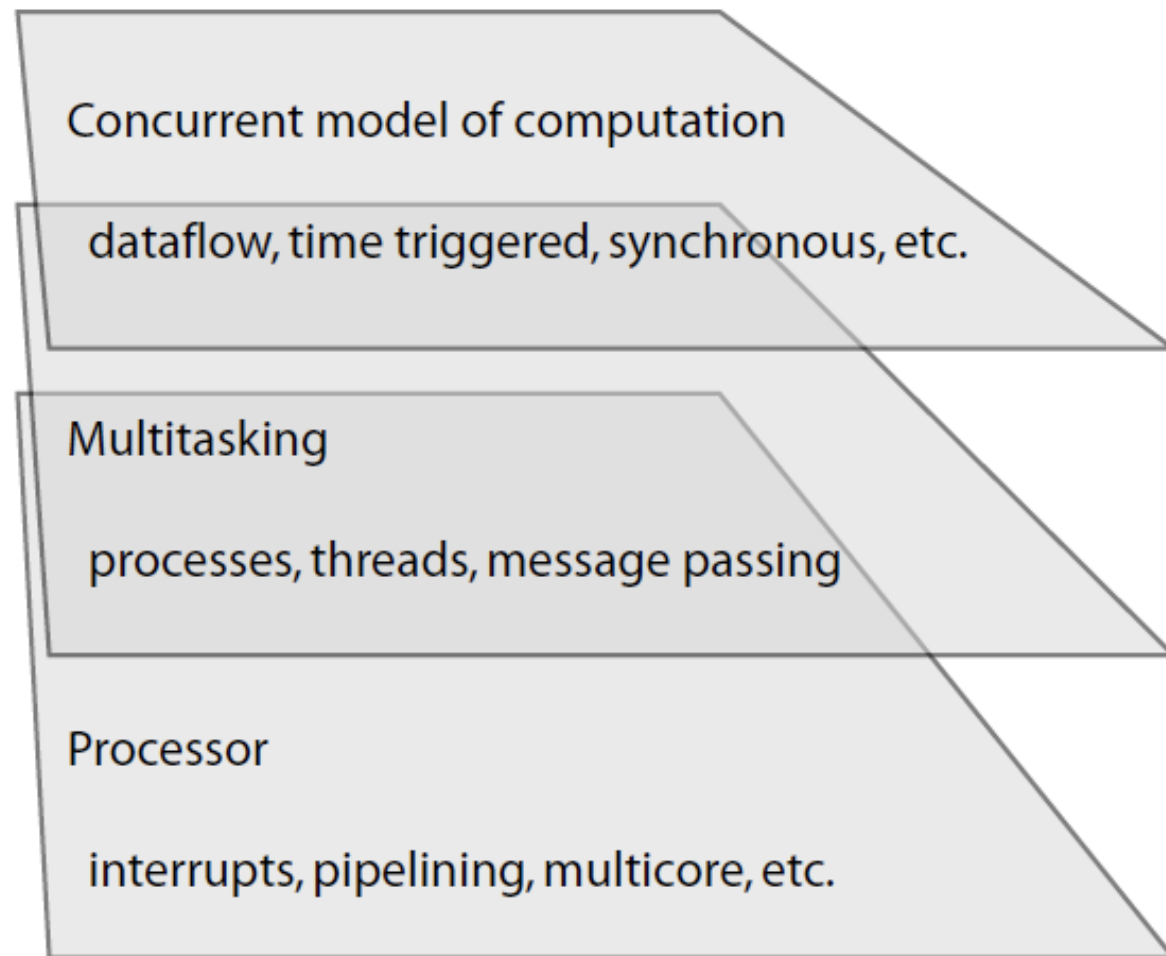
- Threads:
 - Creating/implementing them
 - Memory management
 - Problems associated with threads

11	
Multitasking	
11.1 Imperative Programs	294
<i>Sidebar: Linked Lists in C</i>	<i>297</i>
11.2 Threads	298
11.2.1 Creating Threads	301
11.2.2 Implementing Threads	302
11.2.3 Mutual Exclusion	305
11.2.4 Deadlock	308
<i>Sidebar: Operating Systems</i>	<i>309</i>
11.2.5 Memory Consistency Models	311
11.2.6 The Problem with Threads	316
11.3 Processes and Message Passing	318
11.4 Summary	318
Exercises	318

In this chapter, we discuss mid-level mechanisms that are used in software to provide concurrent execution of sequential code. There are a number of reasons for executing multiple sequential programs concurrently, but they all involve timing. One reason is to improve responsiveness by avoiding situations where long-running programs can block a program that responds to external stimuli, such as sensor data or a user request. Improved responsiveness reduces **latency**, the time between the occurrence of a stimulus and the response. Another reason is to improve performance by allowing a program to run simul-



Layers of Abstraction for Concurrency in Programs



Definition and Uses

- Threads are **sequential procedures** that **share memory**.
- Uses of concurrency:
 - Reacting to external events (interrupts)
 - Exception handling (software interrupts)
 - Creating the illusion of simultaneously running different programs (multitasking)
 - Exploiting parallelism in the hardware (e.g. multicore machines).
 - Dealing with real-time constraints.

Thread Scheduling

Predicting the thread schedule is an **iffy proposition**.

- Without an OS, multithreading is achieved with **interrupts**. Timing is determined by external events.
- Generic OSs (Linux, Windows, OSX, ...) provide **thread libraries** (like “pthreads”) and provide no fixed guarantees about when threads will execute.
- Real-time operating systems (RTOSs), like FreeRTOS, QNX, VxWorks, RTLinux, support a variety of ways of **controlling when threads execute** (priorities, preemption policies, deadlines, ...).
- **Processes** are collections of threads with their own memory, not visible to other processes. **Segmentation faults** are attempts to access memory not allocated to the process. Communication between processes must occur via OS facilities (like pipes or files).

Posix Threads (PThreads)

- **PThreads** is an API (Application Program Interface) implemented by many operating systems, both real-time and not. It is a library of C procedures.
- Standardised by the IEEE in 1988 to unify variants of Unix. Subsequently implemented in most other operating systems.
- An alternative is **Java**, which may use PThreads under the hood, but provides thread constructs as part of the programming language.

Creating and Destroying Threads

```
#include <pthread.h>
```

Can pass in pointers to shared variables.

```
void* threadFunction(void* arg) {  
    ...  
    return pointerToSomething or NULL;  
}
```

Can return pointer to something.

Do not return a pointer to any local variable!

```
int main(void) {  
    pthread_t threadID;  
    void* exitStatus;  
    int value = something;  
    pthread_create(&threadID, NULL, threadFunction, &value);  
    ...  
    pthread_join(threadID, &exitStatus);  
    return 0;  
}
```

Create a thread (may or may not start running!)

Becomes arg parameter to threadFunction.

Why is it OK that this is a local variable?

Return only after all threads have terminated.

What's Wrong with This?

```
#include <pthread.h>
#include <stdio.h>
void *myThread() {
    int ret = 42;
    return &ret;
}
```

Don't return a pointer to a local variable, which is on the stack.

```
int main() {
    pthread_t tid;
    void *status;
    pthread_create(&tid, NULL, myThread, NULL);
    pthread_join(tid, &status);
    printf("%d\n", *(int*)status); return 0;
}
```


Notes

- Threads can (and often do) **share variables**
- Threads may or may not begin running immediately after being created.
- A thread may be **suspended** between any two **atomic instructions** (typically, assembly instructions, not C statements!) to execute another thread and/or interrupt service routine.
- Threads can often be given **priorities**, and these may or may not be respected by the thread scheduler.
- Threads may **block** on semaphores and mutexes (we will do this later in this lecture).


Modeling Threads via Asynchronous Composition of Extended State Machines

States or transitions represent **atomic instructions**


Interleaving semantics:

- Choose one machine, arbitrarily.
- Advance to a next state if guards are satisfied.
- Repeat.

Thread 1



Thread 2



Can Thread 1 be in C1 at the same time Thread 2 is in C2?

Need to **compute reachable states** to reason about correctness of the composed system

A Scenario

Under **Integrated Modular Avionics**, software in the aircraft engine continually runs diagnostics and publishes diagnostic data on the local network.



Proper software engineering practice suggests using the observer pattern.



Design Patterns


Elements of Reusable
Object-Oriented Software

Erich Gamma
Richard Helm
Ralph Johnson
John Vlissides



Foreword by Grady Booch

ADDISON-WESLEY PROFESSIONAL COMPUTING SERIES




An observer process updates the cockpit display based on notifications from the engine diagnostics.

Typical thread programming problem

“The *Observer pattern* defines a **one-to-many dependency** between a subject object and any number of observer objects so that when the subject object changes state, all its observer objects are notified and updated automatically.”

Design Patterns, Eric Gamma, Richard Helm,
Ralph Johnson, John Vlissides
(Addison-Wesley, 1995)



Observer Pattern in C

```
// Value that when updated triggers notification
// of registered listeners.
int value;

// List of listeners. A linked list containing
// pointers to notify procedures.
typedef void* notifyProcedure(int);
struct element {...}
typedef struct element elementType;
elementType* head = 0;
elementType* tail = 0;

// Procedure to add a listener to the list.
void addListener(notifyProcedure listener) {...}

// Procedure to update the value
void update(int newValue) {...}

// Procedure to call when notifying
void print(int newValue) {...}
```

Observer Pattern in C

```
// Value that when updated triggers notification of
// registered listeners.
int value;

// List of listeners. A linked list of pointers to notify procedures.
// pointers to notify procedures.
typedef void* notifyProcedure;
struct element {
    notifyProcedure* listener;
    struct element* next;
};
typedef struct element elementType;
elementType* head = 0;
elementType* tail = 0;

// Procedure to add a listener.
void addListener(notifyProcedure listener) {...}

// Procedure to update the value
void update(int newValue) {...}

// Procedure to call when notifying
void print(int newValue) {...}
```

```
typedef void* notifyProcedure(int);
struct element {
    notifyProcedure* listener;
    struct element* next;
};
typedef struct element elementType;
elementType* head = 0;
elementType* tail = 0;
```

Observer Pattern in C

```
// Value that  
registered lis  
int value;  
  
// List of lis  
// pointers to  
typedef void*  
struct element  
typedef struct  
elementType* h  
elementType* t  
  
// Procedure to  
void addListener  
  
// Procedure to  
void update(in  
  
// Procedure to call when notifying  
void print(int newValue) {...}
```

```
// Procedure to add a listener to the list.  
void addListener(notifyProcedure listener) {  
    if (head == 0) {  
        head = malloc(sizeof(elementType));  
        head->listener = listener;  
        head->next = 0;  
        tail = head;  
    } else {  
        tail->next = malloc(sizeof(elementType));  
        tail = tail->next;  
        tail->listener = listener;  
        tail->next = 0;  
    }  
}
```

Observer Pattern in C

```
// Value that when updated triggers notification of
// registered listeners.
int value;

// List of listeners. A linked list containing
// pointers to notify procedures.
typedef void* notifyProcedure(int);
struct element {...}
typedef struct element* element;
element* head;

// Procedure to update the value
void update(int newValue) {
    value = newValue;
    // Notify listeners.
    element* element = head;
    while (element != 0) {
        (*(element->listener))(newValue);
        element = element->next;
    }
}

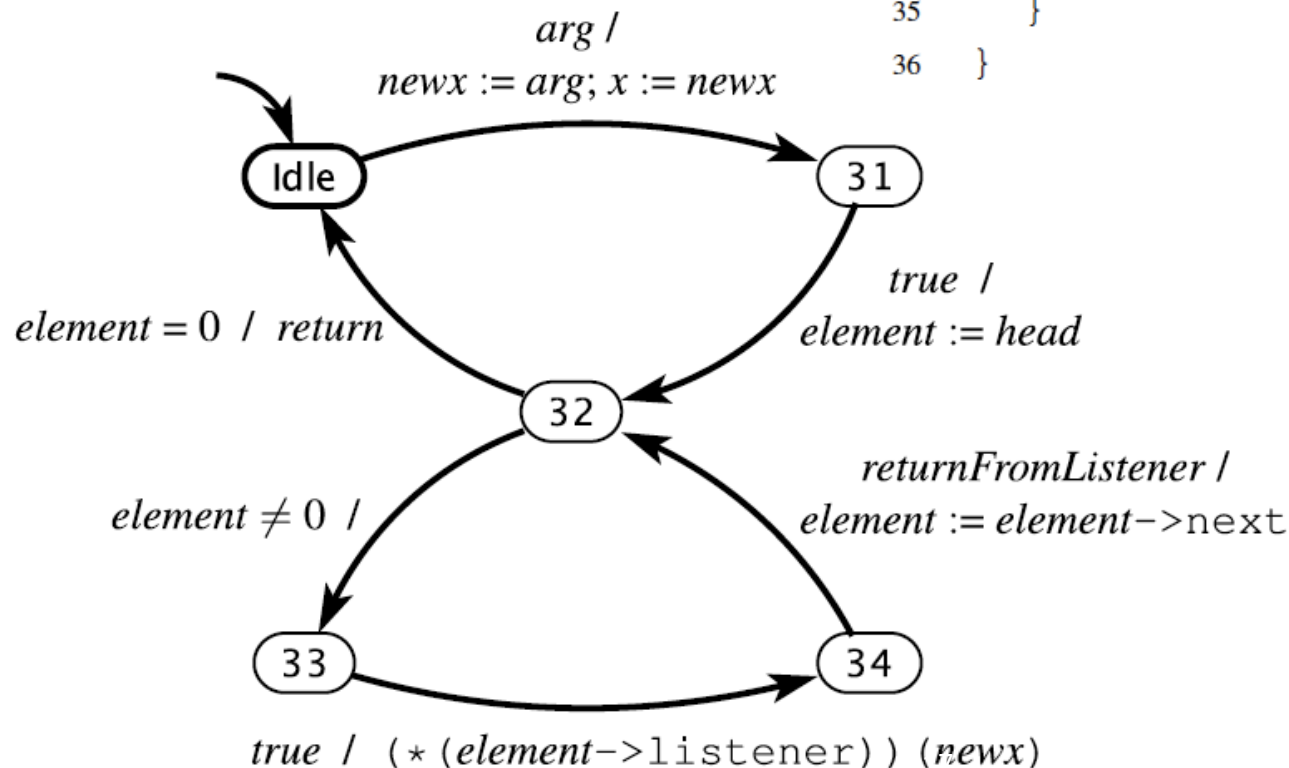
// Procedure to add a listener
void addListener(notifyProcedure* listener) {
    element* element = head;
    while (element != 0) {
        if (element->listener == 0) {
            element->listener = listener;
            return;
        }
        element = element->next;
    }
}

// Procedure to print the value
void print(int value) {
    printf("Value: %d\n", value);
}
```

```
// Procedure to update the value
void update(int newValue) {
    value = newValue;
    // Notify listeners.
    element* element = head;
    while (element != 0) {
        (*(element->listener))(newValue);
        element = element->next;
    }
}
```


Model of the Update Procedure

```
27 // Procedure to update x.  
28 void update(int newx) {  
29     x = newx;  
30     // Notify listeners.  
31     element_t* element = head;  
32     while (element != 0) {  
33         (*(element->listener))(newx);  
34         element = element->next;  
35     }  
36 }
```



Observer Pattern in C

```
// Value that when updated triggers notification of registered listeners.
int value;

// List of listeners. A linked list containing
// pointers to notify procedures.
typedef void* notifyProcedure(int);
struct element {...}
typedef struct element elementType;
elementType* head = 0;
elementType* tail = 0;

// Procedure to add a listener to the list.
void addListener(notifyProcedure listener) {...}

// Procedure to update the value
void update(int newValue) {...}

// Procedure to call when notifying
void print(int newValue) {...}
```

Will this work in a
multithreaded context?

Will there be
unexpected/undesirable
behaviours?

Observer Pattern in C: assume concurrent calls

```
// Value that registered listeners use to pass back their
int value;

// List of listeners
// pointers to listeners
typedef void* listener;
struct element {
    listener listener;
    struct element* next;
};
typedef struct element element;
element* head = 0;
element* tail = 0;

// Procedure to add a listener to the list.
void addListener(notifyProcedure listener) {
    if (head == 0) {
        head = malloc(sizeof(element));
        head->listener = listener;
        head->next = 0;
        tail = head;
    } else {
        tail->next = malloc(sizeof(element));
        tail = tail->next;
        tail->listener = listener;
        tail->next = 0;
    }
}

// Procedure to call when notifying
void update(int newValue) {...}

// Procedure to call when notifying
void print(int newValue) {...}
```

Using Posix mutexes on the observer pattern in C

```
#include <pthread.h>
...
pthread_mutex_t lock;

void addListener(notify listener) {
    pthread_mutex_lock(&lock);
    ...
    pthread_mutex_unlock(&lock);
}

void update(int newValue) {
    pthread_mutex_lock(&lock);
    value = newValue;
    elementType* element = head;
    while (element != 0) {
        (*element->listener)(newValue);
        element = element->next;
    }
    pthread_mutex_unlock(&lock);
}

int main(void) {
    pthread_mutex_init(&lock, NULL);
    ...
}
```

However, this carries a significant **deadlock risk**.

The update procedure **holds the lock** while it calls the notify procedures. If any of those stalls trying to acquire another lock, and the thread holding that lock tries to acquire this lock, **deadlock results**.

```
#include <pthread.h>
...
pthread_mutex_t lock;
```

```
void addListener(notify listener) {
    pthread_mutex_lock(&lock);
    ...
    pthread_mutex_unlock(&lock);
}

void update(int newValue) {
    pthread_mutex_lock(&lock);
    value = newValue;
    ... copy the list of listeners ...
    pthread_mutex_unlock(&lock);
    elementType* element = headCopy;
    while (element != 0) {
        (*(element->listener))(newValue);
        element = element->next;
    }
}

int main(void) {
    pthread_mutex_init(&lock, NULL);
    ...
}
```

One possible “fix”

What is wrong with this?

Notice that if multiple threads call `update()`, the updates will occur in **some order**. But there is no assurance that the listeners will be notified in the same order. Listeners may be misled about the “final” value.

This is a very simple, commonly used design pattern.

Perhaps concurrency is just **hard**...

“Humans are quickly overwhelmed by concurrency and find it much more difficult to reason about concurrent than sequential code. Even careful people miss possible interleavings among even simple collections of partially ordered operations.”

H. Sutter and J. Larus. Software and the concurrency revolution. ACM Queue, 3(7), 2005.

If concurrency were intrinsically hard, we would not function well in the physical world



*It is not
concurrency that
is hard...*

...it is Threads that are hard!

Threads are **sequential processes** that share memory. From the perspective of any thread, the *entire state of the universe can change between any two atomic actions* (itself an ill-defined concept).

Imagine if the physical world did that...

Claim

*Nontrivial software written with threads, semaphores, and mutexes is **incomprehensible** to humans.*

- *Need better ways to program concurrent systems (we will see some later in the course)*
- *Better tools to analyse and reason about concurrency (e.g. model checking)*

Do Threads Have a Sound Foundation?

If the foundation is bad, then we either tolerate brittle designs that are difficult to make work, or we have to rebuild from the foundations.



Succinct Problem Statement

- Threads are wildly **nondeterministic**.
- The programmer's job is to prune away the nondeterminism by **imposing constraints** on execution order (e.g., mutexes) and **limiting shared data accesses** (e.g., OO design).

Incremental Improvements to Threads

- Object Oriented programming
- Coding rules (Acquire locks in the same order...)
- Libraries (Stapl, Java ≥ 5.0 , ...)
- Transactions (Databases, ...)
- Patterns (MapReduce, ...)
- Formal verification (Model checking, ...)
- Enhanced languages (Split-C, Cilk, Guava, ...)
- Enhanced mechanisms (Promises, futures, asynchronous atomic callbacks ...)

Things to do ...

- Read Chapter 12

	12
	Scheduling
12.1 Basics of Scheduling	323
12.1.1 Scheduling Decisions	323
12.1.2 Task Models	325
12.1.3 Comparing Schedulers	327
12.1.4 Implementation of a Scheduler	328
12.2 Rate Monotonic Scheduling	329
12.3 Earliest Deadline First	334
12.3.1 EDF with Precedences	337
12.4 Scheduling and Mutual Exclusion	339
12.4.1 Priority Inversion	339
12.4.2 Priority Inheritance Protocol	340
12.4.3 Priority Ceiling Protocol	342
12.5 Multiprocessor Scheduling	344
12.5.1 Scheduling Anomalies	345
12.6 Summary	348
<i>Sidebar: Further Reading</i>	350
Exercises	351

Chapter 11 has explained **multitasking**, where multiple **imperative** tasks execute concurrently, either interleaved on a single processor or in parallel on multiple processors. When there are fewer processors than tasks (the usual case), or when tasks must be performed at

